



NTNU | Norwegian University of
Science and Technology

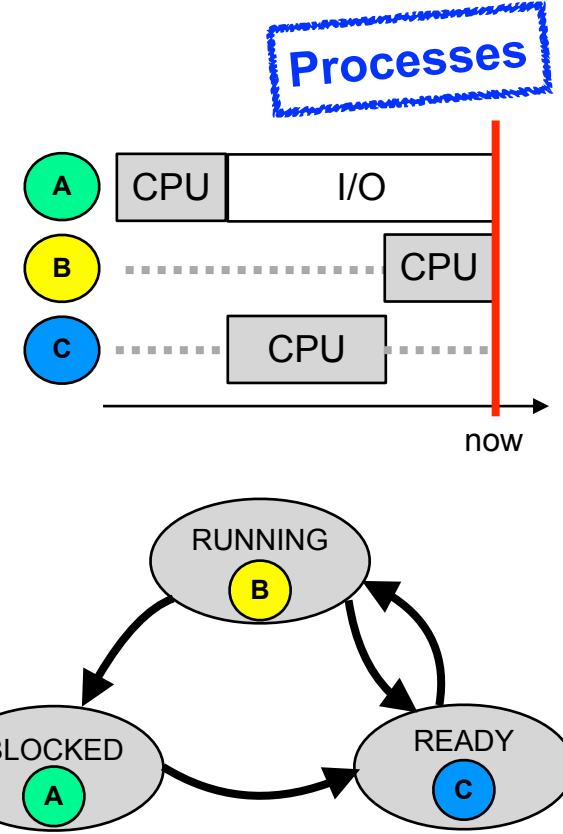
Operating Systems

Lecture 4: Processes

Michael Engel

Review: processes...

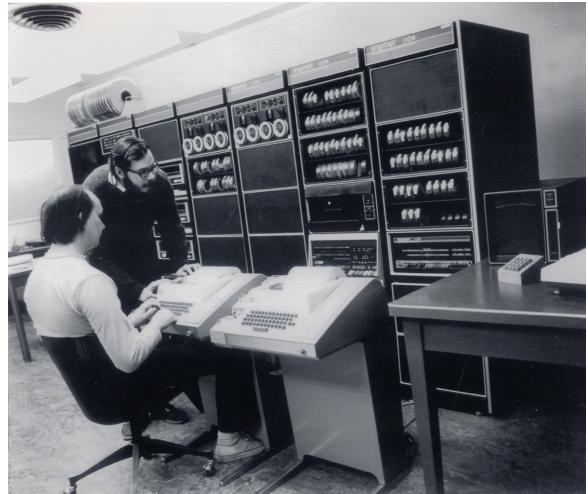
- are “programs in execution”
 - dynamic, not static
 - alternating sequences of “CPU bursts” and “I/O bursts”
- require resources of the computer
 - CPU, memory, I/O devices
- have a state
 - READY, RUNNING, BLOCKED
- are **conceptionally** considered to be independent, concurrent control flows
- are under the control of the operating system
 - resource allocation and revocation



Unix (Thompson & Ritchie 1968)

Unix process model

- A system with a long (hi)story... [1,2]
- Origin: AT&T Bell Labs
 - Developed as an alternative to “Multics”
- Version 1 created on a PDP 7 [4,5]
 - written in assembler
 - 8192 18 bit words of memory
- Version 3 implemented in C on a PDP11
 - C was created to enable OS development in a high-level language



Thompson & Ritchie with a PDP11 at Bell Labs, 1970s

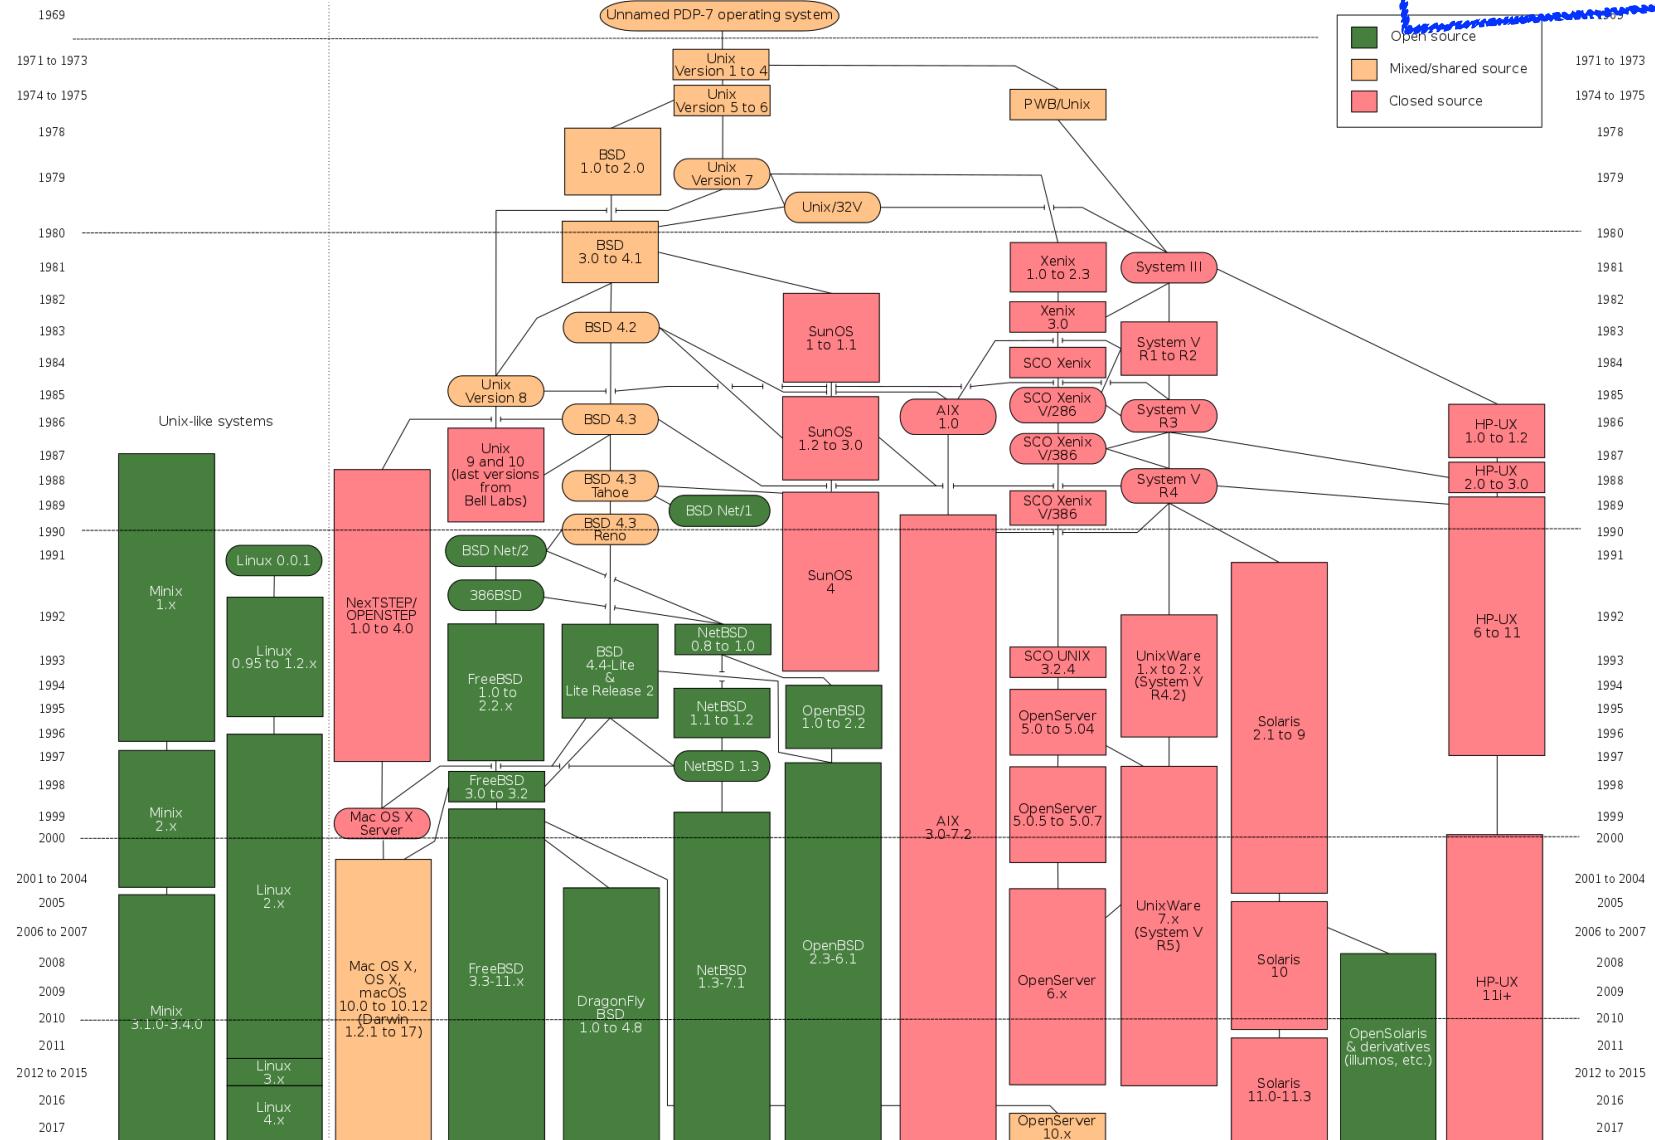


PDP11/40 systems in NTNU's datamuseum

Unix variants

<https://www.levenez.com/unix/>

Unix process model



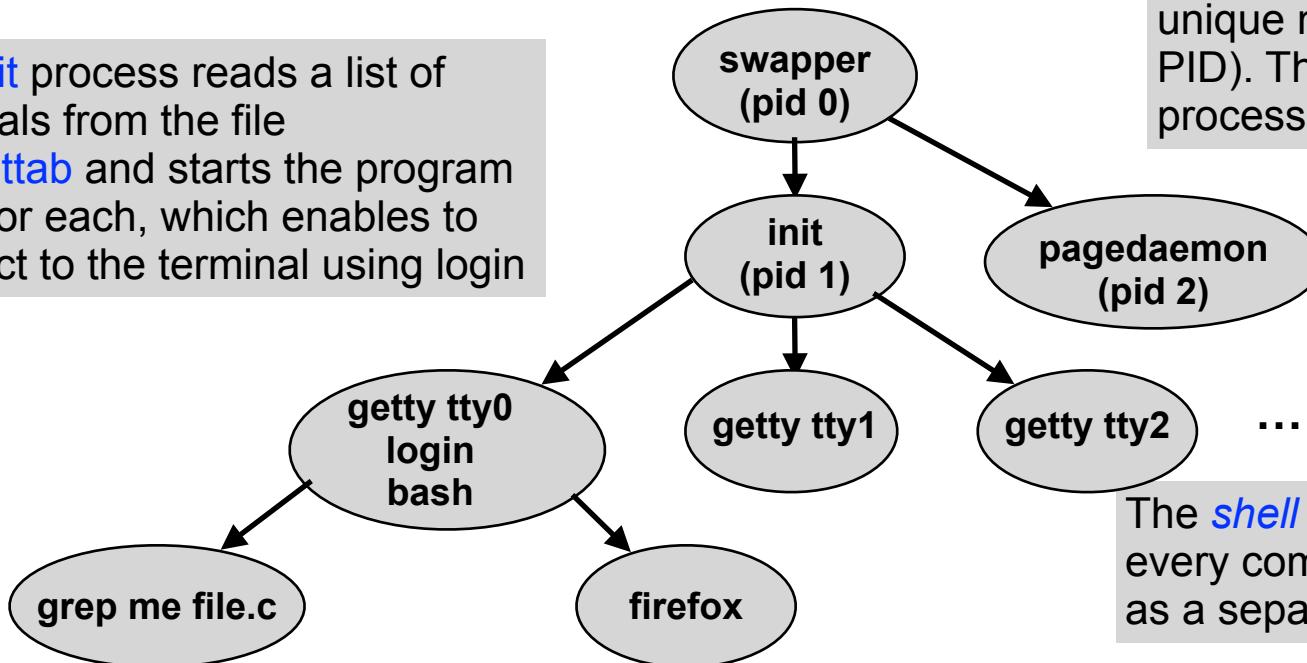
Unix processes

Unix process model

- are the primary way to **structure activities**
 - application as well as system processes
- can create new processes in a fast and easy way
 - parent process → child process
- form a **process hierarchy**:

The `init` process reads a list of terminals from the file `/etc/inittab` and starts the program `getty` for each, which enables to connect to the terminal using login

Every Unix process has a unique number (process id, PID). The **PID** of the parent process is called **PPID**



The **shell** also uses processes: every command is executed as a separate child process

Unix shells

Unix process model

- A “shell” around the operating system “core”
- Text based user interface to start **commands** (Unix programs):
 - Commands can be located anywhere in the file system
 - Shell searches in directories given in the \$PATH **environment variable**
 - e.g. /usr/bin:/bin:/usr/local/bin:....



Shell **prompt**

your **input**

```
me@unix:~> which vim
```

The command **which** shows where a command (here: **vim**) is found in the file system

- Every executed command is a separate **child process**
- Typically, the shell blocks (waits) until the last command has terminated
- It is possible to suspend, continue and terminate commands (*job control*) and to have commands executing in the background

Unix shells: job control

Unix process model



```
me@unix:~> vim foo.c
```

Ctrl-Z

- Command `vim` is started
- The shell blocks

- Command is suspended
- The shell continues to run

```
[1]+ Stopped  vim foo.c
me@unix:~> kate bar.c &
[2] 19504
me@unix:~> jobs
[1]+ Stopped  vim foo.c
[2]- Running  kate bar.c &
me@unix:~> bg %1
[1]+ vim foo.c &
me@unix:~> jobs
[1]- Running  vim foo.c &
[2]+ Running  kate bar.c &
```

- The `&` at the end of the input tells the shell to start and run the `kate` command in the background

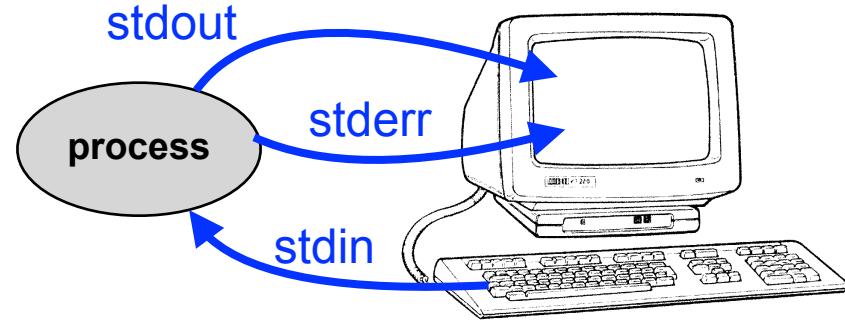
- `jobs` gives a list of all started commands

- `bg` sends a suspended command to the background

Standard I/O channels

Unix process model

- Usually connected to the terminal in which the shell runs that started the process:
 - Standard input (**stdin**): read user input (Keyboard)
 - Standard output (**stdout**): text output of the process (terminal window)
 - Standard error (**stderr**): separate channel for error messages (usually also connected to the terminal)
- Almost all Unix commands also ***accept files as input or output channels*** (instead of the terminal)
- Shells provide a simple syntax to *redirect* the standard I/O channels



Redirecting standard I/O

Unix process model

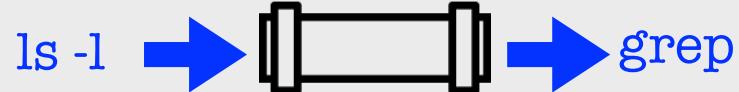
```
me@unix:~> ls -l > d1
me@unix:~> grep "Jan 29" < d1 > d2
me@unix:~> wc < d2
2 18 118
```

Redirection of the **standard output** of the `ls` command to the file `d1` using `>`

Redirection of the **standard input** of the command `wc` from the file `d2` using `<`

This can also be written in a compact way:

```
me@unix:~> ls -l | grep "Jan 29" | wc
2 18 118
```



The **| (pipe)** symbol tells the shell to connect the standard output of the left process (`ls`) to the standard input of the right process (`grep`)

The Unix philosophy

Unix process
model

- Doug McIlroy, the inventor of Unix pipes, summarized the Unix philosophy as follows:

"This is the Unix philosophy:

*Write programs that **do one thing** and do it well.*

*Write programs to **work together**.*

*Write programs to **handle text streams**,
because that is a universal interface."*



This is commonly expressed
in a shorter way:
“Do one thing, do it well.”

Process–OS interaction in Unix

Unix process model

- How does an application program request a service from the operating system?
- From the point of view of the application, calling an operating system service looks like a regular function call, e.g.:

```
pid = fork();
```
- However, arbitrarily calling code inside the OS kernel is dangerous:
 - No checking of permission to execute a function
 - No checking for correct parameters
→ *security nightmare!*
- The transition from code executing in an application to code running in the kernel needs to be protected!

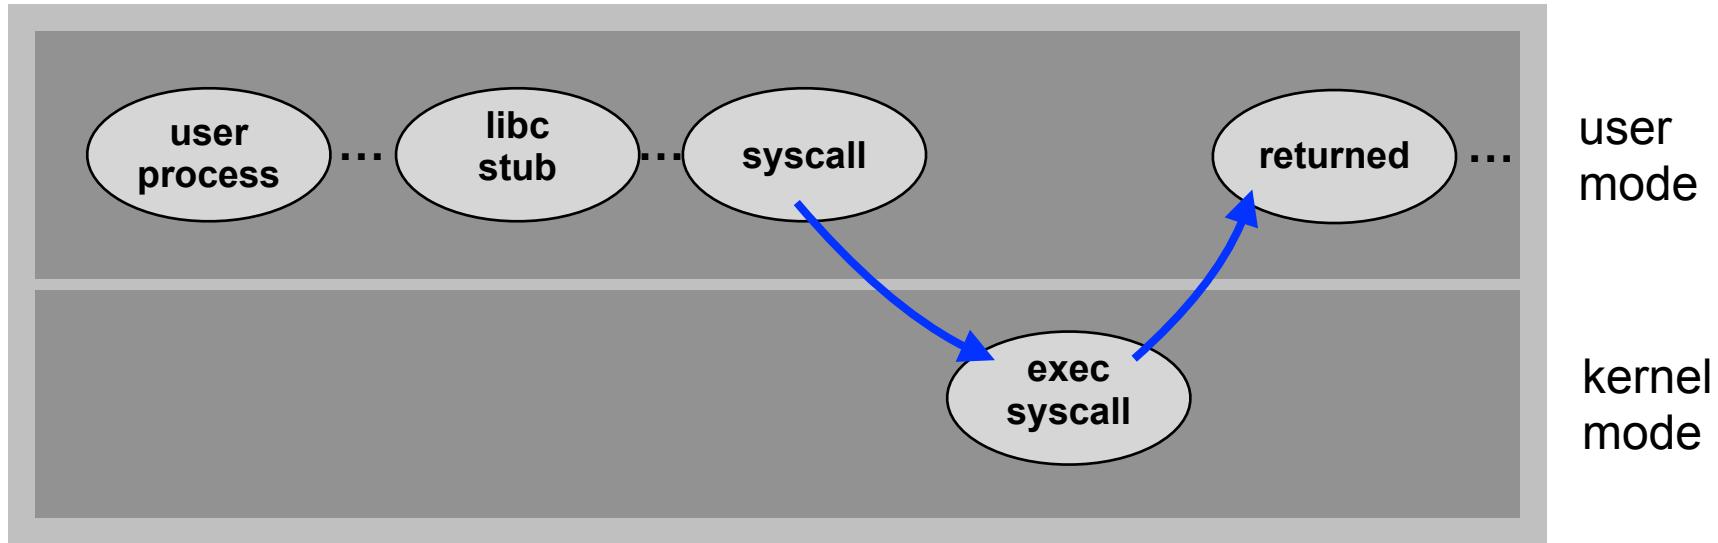
Process–OS interaction in Unix

Unix process model

- The transition from code executing in an application to code running in the kernel needs to be protected!
- Many CPUs provide several *execution modes*:
 - “user mode”: only restricted functionality is allowed
 - “kernel” or “supervisor mode”: full access to all hardware resources
- Special machine instructions are provided to transition from user to kernel mode:
 - int 0x80 (intel x86), syscall/sysenter (intel/AMD64)
 - trap (Motorola 68k), SVC (ARM), ECALL (RISC-V)
- Executing such an instruction causes the CPU to change its current execution mode to kernel mode and jump to an address predetermined by the processor hardware: **system call**

Process–OS interaction in Unix

Unix process model



- Applications can execute a syscall instruction directly, but:
 - This stops working when the syscall interface changes
- In most modern systems, the C library (libc) provides **stubs** (adapter functions) that call the actual syscall
 - The stub function is a regular function linked to the application

Unix process control: syscalls

Unix process model

- A first overview of process related system calls (syscalls) [3]

- `getpid` (2) returns PID of the calling process
- `getppid` (2) returns PID of the parent process(PPID)
- `getuid` (2) returns the UID of the calling process
- `fork` (2) creates a new child process
- `exit` (3), `_exit` (2) terminates the calling process
- `wait` (2) waits for the termination of a child process
- `execve` (2) loads and starts a program in the context of the calling process



The number in brackets gives the **section** of the Unix manual pages the command is described in, read them e.g. with `man 2 wait`

Unix processes in detail: fork()

Unix process model

System call: pid_t `fork` (void)

- Duplicates the calling process
(the standard way to create new processes in Unix!)
- The child process inherits...
 - Address space (code, data, bss, stack segments)
 - User and group ID
 - Standard I/O channels
 - Process group, signal table (more on this later)
 - Open files, current working directory (also later...)
- Not** copied are the following:
 - Process ID (PID), parent process ID (PPID)
 - Pending signals, accounting data, ...
- One process calls fork, but two processes return**

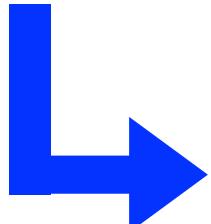
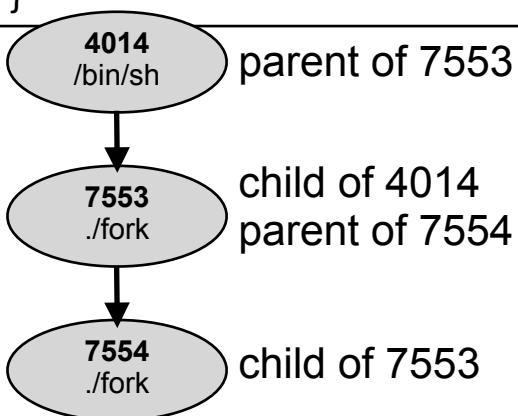
This is the C prototype
for fork:
no parameters (void),
returns a pid_t value



Use of fork()

Unix process model

```
... /* includes */  
int main () {  
    int pid;  
    printf("In parent: pid %d PPID %d\n", getpid(), getppid());  
    pid = fork(); /* Process is duplicated here!  
                    Both continue running from here */  
    if (pid > 0)  
        printf("In the parent process, child PID %d\n", pid);  
    else if (pid == 0)  
        printf("In the child process, PID %d PPID %d\n",  
               getpid(), getppid());  
    else  
        printf("Oh, an error!\n"); /* more in the theoretical ex.* */  
}
```



```
me@unix:~> gcc -o fork fork.c  
me@unix:~> ./fork  
In parent: pid 7553 PPID 4014  
In the child process, PID 7554 PPID 7553  
In the parent process, chip PID 7554
```

Discussion: fast process creation

Unix process model

- Copying the address space takes a lot of time
 - Especially if the program immediately calls `exec..()` afterwards
→ *complete waste of time!*
- Historic solution: `vfork`
 - The parent process is suspended until the child process calls `exec..()` or terminates using `_exit()`
- The child simply uses code and data of its parent (without copying!)
 - The child process ***must not change any data***
 - sometimes not so simple: e.g., don't call `exit()`, but `_exit()`!
- Modern solution: `copy on write`
 - Parent and child process ***share the same code and data segments*** using the memory management unit (MMU)
 - A segment is copied only if the child process changes any data
 - This is not the case when `exec..()` is called directly after `fork()`
 - `fork()` using copy on write is *almost as fast as vfork()*

Unix processes in detail: `_exit()`

Unix process model

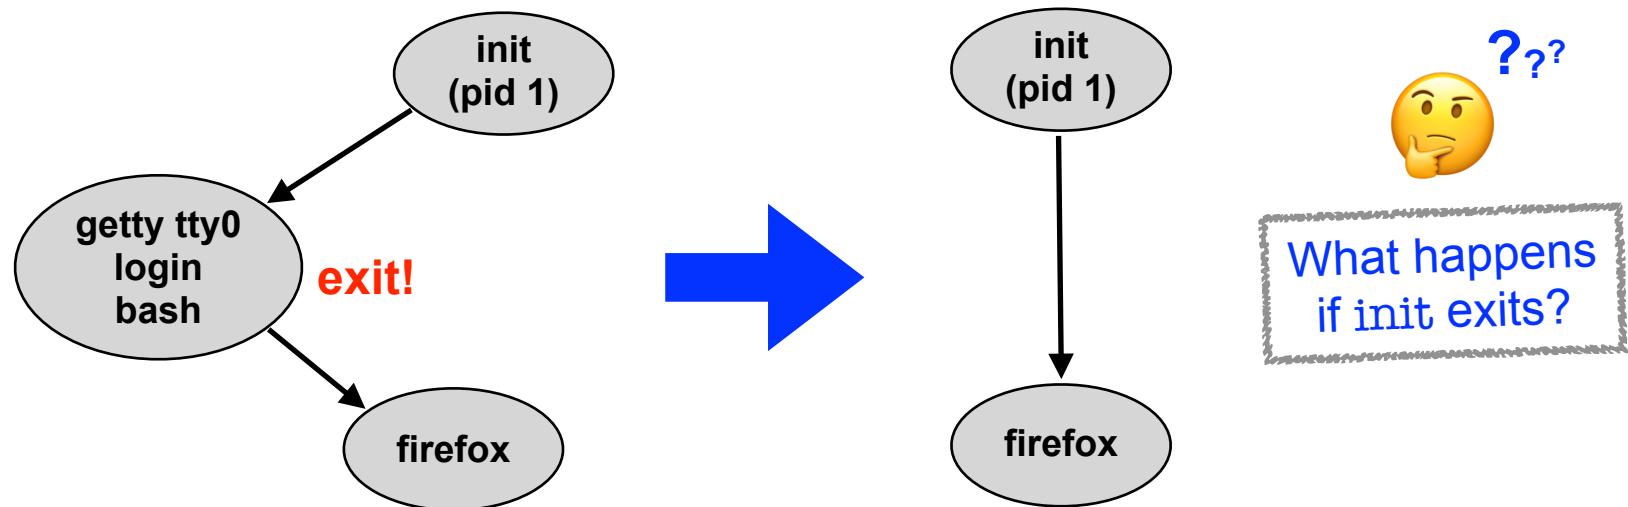
System call: void `_exit(int)`

- Terminates the calling process and passes the integer argument as “exit status” to the parent process
 - This call **does not return!**
- Releases the resources allocated by the process
 - open files, used memory, ...
- Sends a signal SIGCHLD to its parent process
- There is also a library function `exit(3)` which additionally releases resources used by the C library
 - Among other things, this outputs (*flushes*) all data still stored in output buffers!
- Normal processes should use `exit(3)`, not `_exit`

Discussion: orphaned processes

Unix process model

- A Unix process is **orphaned** when its parent process terminates
- What happens to our process hierarchy?



The **init** process (always pid 1) adopts all orphaned processes. Thus, the process hierarchy is still in working order.

Unix processes in detail: `wait()`

Unix process model

System call: `pid_t wait(int *)`

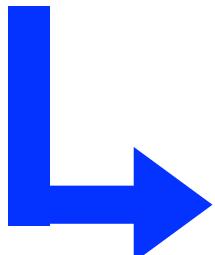
- Blocks the calling process until one of its child processes terminates
 - The return value of `wait` is the terminated child's PID
 - Using the `int *` parameter, the caller is passed the child's “exit status” (and more)
- `wait` returns immediately if all child processes are already terminated

Use of `wait()`

Unix process model

```
... /* includes, main() { ... */  
pid = fork(); /* create child process */  
if (pid > 0) {  
    int status;  
    sleep(5); /* Library function: sleep 5 seconds */  
    if (wait(&status) == pid && WIFEXITED(status))  
        printf ("Exit status: %d\n", WEXITSTATUS(status));  
}  
else if (pid == 0) {  
    exit(42);  
}  
...
```

A process can also be “**killed**” from the outside, i.e. it does not call `exit`. In this case, `WIFEXITED` would return 0.



```
me@unix:~> ./wait  
Exit status: 42
```

Discussion: zombies

Unix process model

- A terminated process is called a “**zombie**” until its exit status is requested using `wait`
- The resources allocated to such processes can be released, but the OS process management still needs to know about them
- Especially the *exit status* has to be saved



[by Charlie Llewellyn, CC BY 2.0]

```
me@unix:~> ./wait & ←  
me@unix:~> ps  
 PID TTY      TIME CMD  
 4014 pts/4    00:00:00 bash  
17892 pts/4    00:00:00 wait  
17895 pts/4    00:00:00 wait <defunct>  
17897 pts/4    00:00:00 ps  
me@unix:~> Exit status: 42
```

Example program from the previous slide during the 5 seconds waiting time

Zombies are annotated by `ps` as `<defunct>`

Unix processes in detail: execve()

Unix process model

System Call: `int execve (const char *command,
const char *args[], const char *envp[])`

- Loads and starts the command passed in the “command” parameter
- Only returns in case of an error
 - e.g. command does not exist, no access, ...
- Replaces the complete address space of the calling process
 - ***but it remains the same process!***
 - Same PID, PPID, open files, ...
- The C library (libc) provides some comfortable support functions that internally call `execve`:
`execl`, `execv`, `execlp`, `execvp`, ... (check their man pages!)

Use of exec()

Unix process model

```
... /* includes, main() { ... */  
char cmd[100], arg[100];  
while (1) {  
    printf ("Command?\n");  
    scanf ("%99s %99s", cmd, arg);  
    pid = fork(); /* Process is duplicated!  
                    Both continue running from here. */  
    if (pid > 0) {  
        int status;  
        if (wait(&status) == pid && WIFEXITED(status))  
            printf ("Exit Status: %d\n", WEXITSTATUS(status));  
    }  
    else if (pid == 0) {  
        execlp(cmd, cmd, arg, NULL);  
        printf ("exec failed\n");  
    }  
    ...  
}
```

Discussion: why no forkexec()

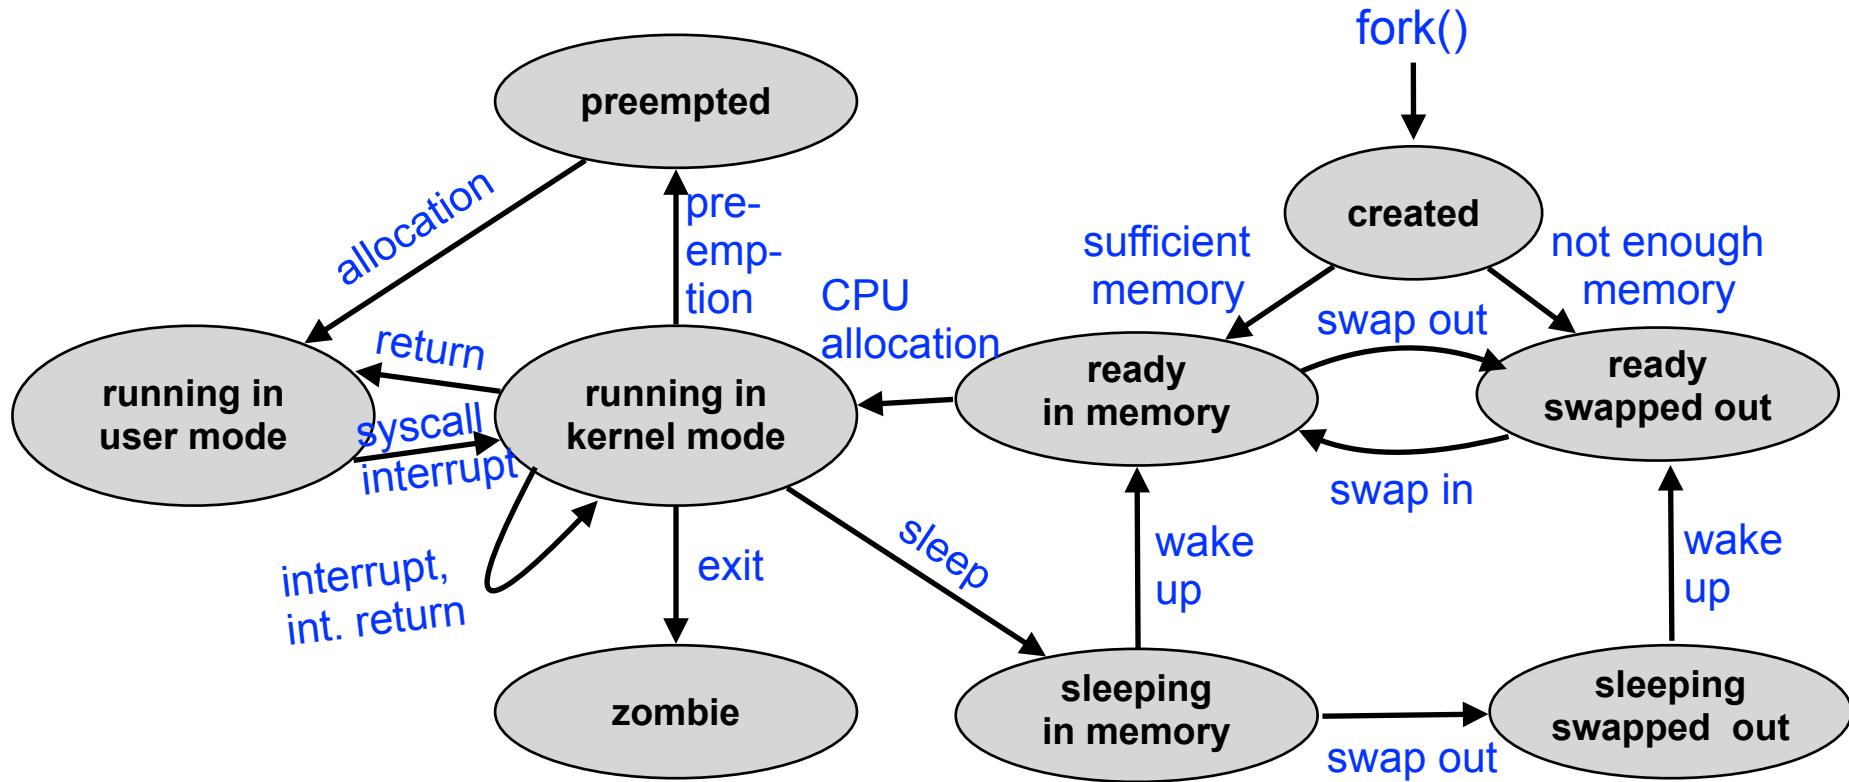
Unix process model

- The parent process has more control if we separate the calls to `fork` and `execve`:
 - Execute operations in the context of the child process
 - Full access to the parent processes data
- Unix shells use this feature to e.g.
 - redirect the standard I/O channels
 - configure pipes

Unix process states

Unix process model

- a bit more complex than our earlier simple model...



Conclusion

- Process management is an important part of any OS
 - Unix has a *process hierarchy*
 - The *init* process (PID 1) is the root of the hierarchy
- Special approach taken in Unix:
separate process creation (`fork`) and program execution (`exec`)!
 - Used by the Unix shell to implement I/O redirection
- Small set of basic system calls for process management
 - Hardware support required to make `fork` efficient
- Real-world process states are quite complex

References

1. Peter H. Salus, A Quarter Century of Unix, Addison-Wesley 1995, ISBN-13: 978-0201547771
2. Brian Kernighan, UNIX: A History and a Memoir, Independently published 2019, ISBN-13: 978-1695978553
3. W. Stevens, Stephen Rago, Advanced Programming in the UNIX Environment, 3rd Edition, Addison-Wesley 2013, ISBN-13: 978-0321637734
4. Dennis M. Ritchie and Ken Thompson. 1974. The UNIX time-sharing system. Commun. ACM 17, 7 (July 1974), 365–375.
DOI:<https://doi.org/10.1145/361011.361061>
5. Dennis M. Ritchie, The Unix Time-Sharing System - A Retrospective, <https://www.bell-labs.com/usr/dmr/www/retro.pdf>